



SI-Special Issue on Security and Privacy in Federated Learning and Unlearning

SecFDCC: Navigating the Federated Clustering Tradeoff with Spectral Embedding and Homomorphic Aggregation

Submission ID82af64e8-bd93-48cc-8c0e-c9084e5e0b1c

Submission VersionInitial Submission

PDF Generation10 Mar 2026 06:08:52 EST by Atypon ReX

Authors

Prof. Lin Lu
Corresponding Author
Submitting Author

 [ORCID](https://orcid.org/0000-0002-2803-6297)
<https://orcid.org/0000-0002-2803-6297>

- Affiliations**
- the Hubei Key Laboratory of Intelligent Vision Based Monitoring for Hydroelectric Engineering
 - the College of Computer and Information Technology, China Three Gorges University, Yichang, 443002, China

Mr. Jiaxin Lin

- Affiliations**
- the Hubei Key Laboratory of Intelligent Vision Based Monitoring for Hydroelectric Engineering

Mr. Rodney Donald Hounkanrin

- Affiliations**
- the Hubei Key Laboratory of Intelligent Vision Based Monitoring for Hydroelectric Engineering

Prof. Jianting Ning

 [ORCID](https://orcid.org/0000-0001-7165-398X)
<https://orcid.org/0000-0001-7165-398X>

- Affiliations**
- the School of Cyber Science and Engineering, Wuhan University, Wuhan 430072, China
 - the Faculty of Data Science, City University of Macau, Macau, China
 - the Hubei Key Laboratory of Intelligent Vision Based Monitoring for Hydroelectric Engineering

Additional Information

Keywords

H Information Technology and Systems / H.2 Database Management / H.2.8 Database Applications / H.2.8.b Clustering, classification, and association rules

H.2.8.b Clustering, classification, and association rules

I Computing Methodologies / I.2 Artificial Intelligence / I.2.6 Learning / I.2.6.g Machine learning

I.2.6.g Machine learning

I Computing Methodologies / I.5 Pattern Recognition / I.5.3 Clustering

I.5.3 Clustering

Files for peer review

All files submitted by the author for peer review are listed below. Files that could not be converted to PDF are indicated; reviewers are able to access them online.

Name	Type of File	Size	Page
SecFDCC_Main.pdf	Main Document - PDF	700.9 KB	Page 4
SecFDCC_Supp.pdf	Appendices	2.4 MB	Page 16

SecFDCC: Navigating the Federated Clustering Tradeoff with Spectral Embedding and Homomorphic Aggregation

Lin Lu, *Member, IEEE*, Jiaxin Lin, Rodney Donald Hounkanrin, Jianting Ning, *Senior Member, IEEE*

Abstract—Federated Clustering has emerged as a vital technique for discovering latent structures in decentralized unlabeled data while preserving privacy. However, real-world deployments face a critical “tradeoff triangle” involving non-IID data robustness, rigorous privacy protection, and resource efficiency. Existing methods typically sacrifice one attribute to optimize the others. In this paper, we propose SECFDCC (Secure Federated Dual-Space Complementary Clustering), a novel framework that navigates this tradeoff. SecFDCC introduces a dual-space mechanism that aligns local cluster structures in both the original feature space and a spectral embedding space via global anchors, effectively handling severe label skew and feature shifts. To guarantee privacy without utility loss, we design a secure aggregation protocol based on CKKS homomorphic encryption, ensuring zero leakage of intermediate gradients or cluster centers. Furthermore, we optimize communication scalability by transmitting compact anchor-based statistics, making bandwidth usage independent of local dataset sizes. Extensive experiments on benchmark datasets demonstrate that SecFDCC outperforms state-of-the-art baselines by 4-12% in accuracy while maintaining optimal privacy resistance against membership inference attacks. The framework is particularly well-suited for privacy-sensitive domains like healthcare and finance where regulatory compliance and data protection are paramount.

Index Terms—Federated Learning, Federated Clustering, Homomorphic Encryption, Spectral Clustering, Secure Computation

I. INTRODUCTION

Federated learning (FL) traditionally focuses on supervised tasks [1, 2], yet acquiring high-quality distributed labels is often expensive, infeasible, or privacy-invasive in domains like healthcare and IoT. This has driven a vital shift toward unsupervised *federated clustering* (FC). However, realizing effective FC in real-world deployments faces two critical, intertwined obstacles:

- 1) **Non-IID Data Heterogeneity**: Severe variations in local data distributions (e.g., label skew, feature shift, quantity imbalance) cause local clusters to diverge from the global structure, making naive aggregation unstable and ineffective.

Manuscript received xxx; revised xxx; accepted xxx. This work was supported in part by the CTGU Advanced Computing Center.

Jianting Ning is with the School of Cyber Science and Engineering, Wuhan University, Wuhan 430072, China, and also with the Faculty of Data Science, City University of Macau, Macau, China. Other authors are with the Hubei Key Laboratory of Intelligent Vision Based Monitoring for Hydroelectric Engineering and also the College of Computer and Information Technology, China Three Gorges University, Yichang, 443002, China (Corresponding author: Lin Lu, e-mail: lulin@ctgu.edu.cn).

- 2) **Privacy Leakage**: Transmitting plaintext cluster statistics or sample assignments exposes sensitive client data to reconstruction and membership inference attacks [3, 4]. Strong cryptographic mechanisms (e.g., homomorphic encryption) are essential for highly regulated domains.

Existing methods typically compromise one aspect for the other: robust methods often leak plaintext statistics, while privacy-preserving approaches degrade under severe heterogeneity or rely on lossy differential privacy. To bridge this critical gap, we propose SECFDCC (Secure Federated Dual-Space Complementary Clustering), a framework that simultaneously achieves rigorous privacy protection and robust clustering under severe non-IID conditions. Our core contributions are threefold:

- **Dual-Space Robustness against Non-IID Data**: SecFDCC jointly optimizes cluster compactness in the original feature space and structural consistency in a spectral embedding space. By aligning local topologies via shared global anchors, it effectively corrects biases caused by label skew and feature shifts.
- **Cryptographically-Sound Aggregation**: Utilizing CKKS homomorphic encryption (HE), the server updates global centers and aligns embeddings entirely over ciphertexts. This guarantees *zero utility loss* compared to plaintext baselines while providing semantic security against honest-but-curious servers and inference attacks.
- **Anchor-Based Communication Efficiency**: By encoding client contributions through a small set of anchor-related statistics rather than expansive affinity matrices or per-sample assignments, SecFDCC bounds communication complexity to $O(K)$ per client. This overhead is strictly **independent of local dataset sizes**, ensuring practical scalability.

The rest of this paper is organized as follows: Section II reviews related work. Section III formulates the problem and threat model. Section IV details the SECFDCC framework. Section VI presents comprehensive evaluations, and Section VII concludes.

II. RELATED WORK

Federated learning (FL) enables collaborative model training without sharing raw data. While supervised FL is extensively studied, *federated clustering* (FC) for discovering

latent structures in distributed unlabeled data remains under-explored. Early FC research primarily adapted K-means and fuzzy clustering to distributed settings [5, 6, 7, 8, 9, 10]. However, these iterative methods typically assume moderately balanced data and are highly vulnerable to non-IID distributions and privacy leakage, motivating the need for secure and robust designs.

A. Federated Clustering under Non-IID Data

Real-world FL systems inevitably encounter highly non-IID data [11], where label skew and feature shifts cause naive centroid averaging to yield biased global models. To mitigate this, one-shot FC methods [12, 13, 14] attempt to reduce communication by collecting local centers or density cores just once. While efficient, these methods remain vulnerable when local cluster structures are severely distorted by extreme heterogeneity.

Recent advances in federated spectral clustering have shown promise in handling heterogeneous data. Qiao et al. [15] and Thakkar and Joshi [16] utilized spectral embeddings to capture intrinsic manifold structures without accessing raw data. However, these methods primarily operate in a single space (graph-based or spectral) and focus on similarity matrix reconstruction, lacking mechanisms for dual-space integration needed to preserve interpretability for non-graph tabular data. Inspired by these findings, our method introduces a *dual-space complementary enhancement* mechanism that jointly leverages the original feature space and a spectral embedding space, utilizing anchor-augmented graphs to maintain alignment under severe non-IID settings.

B. Privacy-Preserving Federated Clustering

Sharing model updates or clustering statistics can expose sensitive information to reconstruction and membership inference attacks [4, 17, 18], necessitating joint consideration of privacy and robustness [19, 20, 21, 22]. Early privacy-preserving FC focused on securing simple statistics (e.g., per-cluster counts and feature sums) using elliptic-curve encryption or secret sharing [23, 24, 25, 26, 27].

To strengthen privacy, recent approaches combine secure multiparty computation (SMC) with Differential Privacy (DP) [28, 29]. In FC, Laplace noise is often added to protect distance and density-core statistics [13, 14]. While DP-based approaches provide formal privacy guarantees, they inevitably face a severe accuracy-privacy trade-off, which degrades clustering utility, particularly for high-dimensional features.

C. Homomorphic Encryption and Secure Aggregation

Homomorphic Encryption (HE) has become a cornerstone of secure aggregation in FL, allowing computations directly on ciphertexts [30]. Practical secure aggregation protocols [31, 32] enable servers to aggregate high-dimensional client updates without learning individual contributions. However, while HE is widely used in supervised FL, FC literature still largely relies on weaker encryption schemes [24, 25] that only




Fig. 1: The common forms of cluster heterogeneity.

protect low-order statistics and require multiple communication rounds, amplifying privacy risks and burdening resource-constrained clients [33].

In summary, existing FC methods suffer from intertwined limitations: they lack the ability to simultaneously address severe data heterogeneity and strict privacy preservation without incurring significant utility loss. SecFDCC fills this gap by providing a robust, dual-space FC mechanism secured end-to-end by advanced HE, ensuring zero utility degradation in adversarial and highly heterogeneous environments.

III. PROBLEM FORMULATION AND CHALLENGES

Consider an unsupervised FL scenario with M clients, where each client $i \in \{1, 2, \dots, M\}$ possesses a private dataset $X_i = \{\mathbf{x}_i^{(1)}, \mathbf{x}_i^{(2)}, \dots, \mathbf{x}_i^{(n_i)}\}$ with $\mathbf{x}_i^{(j)} \in \mathbb{R}^d$. The objective is to discover K global cluster centers $\mathcal{C} = \{\mathbf{c}_1, \mathbf{c}_2, \dots, \mathbf{c}_K\}$ that minimize the global clustering objective:

$$\min_{\mathcal{C}} \sum_{i=1}^M \sum_{j=1}^{n_i} \min_{k \in [K]} \|\mathbf{x}_i^{(j)} - \mathbf{c}_k\|^2, \quad (1)$$

subject to constraints elaborated in the following subsections.

A. Cluster Heterogeneity

In real-world federated scenarios, data across clients often exhibit severe non-IID characteristics, leading to the cluster heterogeneity. This heterogeneity manifests in three forms (Figure 1):

- 1) **Label Distribution Skew:** Client i may have samples from only a subset of clusters (e.g., medical data segregated by hospital specialty);
- 2) **Feature Distribution Shift:** Same cluster exhibits different feature distributions across clients (e.g., regional language variations);
- 3) **Quantity Imbalance:** Severe imbalance in cluster sizes across clients (e.g., rare disease detection).

Under cluster heterogeneity, three main challenges arise: (i) missing clusters at local clients cause center drift in vanilla FedAvg-based clustering; (ii) Euclidean distance in the original feature space fails to capture manifold structures under distribution shifts; and (iii) simple averaging of cluster centers biases models toward clients with larger data volumes.

B. Threat Model and Privacy Risks

We consider a standard honest-but-curious threat model [2, 19], where the central server faithfully executes the protocol




Fig. 2: Overview of our SecFDCC working flow

but may try to infer client-side private information from all received messages. An external eavesdropper may also passively intercept communications. Clients themselves are assumed not to collude. Under this setting, directly sharing local clustering outputs in FC is highly sensitive: local centroids, affinity matrices, or anchor points all encode rich structural information about the underlying data. From these objects, an adversarial server can approximate individual data points, infer cluster memberships and label distributions, or reconstruct local manifolds, leading to serious privacy leakage beyond what is typically seen in supervised FL.

IV. SECFDCC METHOD

Figure 2 illustrates our *SecFDCC* framework for privacy-preserving, robust FC. The server maintains synchronized cluster centers in both the original data space and a spectral embedding space, then each communication round proceeds as follows:

The server broadcasts these dual centers to all participating clients. Each client augments its local dataset with the global centers (used as anchors) to construct an anchor-augmented graph, derives a normalized spectral embedding, and aligns it to the global coordinate frame via orthogonal Procrustes analysis—addressing potential inconsistencies caused by Non-IID heterogeneity. Using both original-space and embedding-space centers, each client computes dual-space fuzzy memberships and related local statistics and encrypts the results using homomorphic encryption. Clients then transmit these encrypted statistics to the server, which securely aggregates the data and updates the global centers using a topology consistency optimization.

Algorithm 1 details the step-by-step protocol, with non-IID robustness procedures such as spectral embedding, dual-space membership calculation, and topology consistency enforcement emphasized in blue, and privacy-protecting mechanisms

(including CKKS encryption and secure aggregation) highlighted in red. The algorithm clarifies server- and client-side roles, cleanly separates the anchor mechanism, and is faithful to practical implementation.

A. Dual-Space Cluster Centers

To address non-IID challenges, SecFDCC maintains two sets of cluster centers:

- **Original space:** $\mathcal{C}_{\text{orig}} = \{\mathbf{c}_k \in \mathbb{R}^d\}_{k=1}^K$, directly interpretable and applicable to new data in the feature space;
- **Embedding space:** $\mathcal{C}_{\text{emb}} = \{\mathbf{e}_k \in \mathbb{R}^{d' \times d}\}_{k=1}^K$, representing clusters in a spectral embedding derived from pairwise affinities.

Maintaining cluster centers in the spectral embedding space serves two purposes. First, spectral embeddings filter out noise by projecting data onto the leading eigenvectors of the graph Laplacian, retaining only the most informative pairwise relationships; centers defined in this space are therefore more robust to irrelevant feature variations. Second, unlike Euclidean distance in the original space, spectral embeddings can capture complex, nonlinear manifold structures through pairwise affinities, making non-convex clusters linearly separable. This is especially valuable under Non-IID heterogeneity, where each client observes only a partial view of the global manifold. The embedding-space centers thus act as global anchors that help align local structures across clients, complementing the original-space centers which preserve direct feature-level interpretability.

The dual-space clustering extends the objective Eq. 1 with balancing cluster quality of in both spaces:

$$\alpha \sum_{i,j,k} u_{ijk}^m \|\mathbf{x}_i^{(j)} - \mathbf{c}_k\|^2 + (1 - \alpha) \sum_{i,j,k} u_{ijk}^m \|\hat{\mathbf{h}}_i^{(j)} - \mathbf{e}_k\|^2 \quad (2)$$

where u_{ijk} is the fuzzy membership of point j at client i to cluster k , $\hat{\mathbf{h}}_i^{(j)}$ is the j -th row of embedding matrix $\hat{\mathbf{H}}_i$

Algorithm 1: Secure Federated Dual-Space Complementary Clustering (SecFDCC)

Require: Clients $\{1, \dots, M\}$ with non-IID data $\{X_i\}$, number of clusters K , embedding dimension d' , total rounds T , fuzziness m , fusion coefficient α

Ensure: CKKS key pair (pk, sk) ; Global cluster centers $C_{\text{orig}} \in \mathbb{R}^{K \times d}$, $C_{\text{emb}} \in \mathbb{R}^{K \times d'}$

- 1: **Server:** Generate CKKS key pair (pk, sk)
- 2: **Server:** Initialize C_{orig} , C_{emb} (e.g., random/k-means++)
- 3: **for** round $t = 1$ to T **do**
- 4: **Server:** Broadcast current C_{orig} , C_{emb} , pk to selected clients
- 5: **for** each client i **in parallel do**
- 6: // **Anchor-based embedding and alignment**
- 7: Form extended data: $X_i^{\text{ext}} = [X_i; C_{\text{orig}}]$
- 8: Build local affinity matrix; compute normalized Laplacian L_i^{ext}
- 9: Spectral embedding: $H_i^{\text{ext}} \leftarrow$ eigenvectors of L_i^{ext} for d' smallest eigenvalues
- 10: Partition: $H_i \leftarrow$ first n_i rows, $H_{\text{anchor}}^{(i)} \leftarrow$ last K rows of H_i^{ext}
- 11: Procrustes alignment: Find R_i^* to align $H_{\text{anchor}}^{(i)}$ to C_{emb} ; set $\hat{H}_i = H_i R_i^*$
- 12: // **Dual-space fuzzy membership and sufficient statistics**
- 13: Compute cluster memberships in original: $U_{\text{orig}} \leftarrow \text{fuzzy}(X_i, C_{\text{orig}})$ and in aligned embedding: $U_{\text{emb}} \leftarrow \text{fuzzy}(\hat{H}_i, C_{\text{emb}})$
- 14: Fuse: $U_{\text{fused}} = \alpha U_{\text{orig}} + (1 - \alpha) U_{\text{emb}}$
- 15: Compute $U_m = U_{\text{fused}}^m$ (elementwise power)
- 16: Sufficient statistics: $S_{\text{orig}} = U_m^\top X_i$, $S_{\text{emb}} = U_m^\top \hat{H}_i$, $n_k = \text{sum}(U_m, \text{rows})$
- 17: **Encrypt** $(S_{\text{orig}}, S_{\text{emb}}, n_k)$ with pk and upload to server
- 18: **end for**
- 19: **Server:** Aggregate encrypted statistics from all clients (homomorphic addition)
- 20: **Server:** Decrypt aggregate stats and update $C_{\text{orig}}, C_{\text{emb}}$ (homomorphic multiplication)
- 21: **Server:** Refine $C_{\text{orig}}, C_{\text{emb}}$ by optimizing topology consistency loss
- 22: Check convergence; **break** if criteria met
- 23: **end for**
- 24: **return** $C_{\text{orig}}, C_{\text{emb}}$

(Eq. (9)), $m > 1$ is the fuzziness parameter (typically $m = 2$), and $\alpha \in [0, 1]$ balances the two spaces (typically $\alpha = 0.5$).

Theorem 1 (Dual-Space Non-IID Robustness). *Under non-IID data with label distribution skew, the dual-space objective (2) achieves clustering error:*

$$\mathcal{E}_{\text{dual}} \leq \alpha \cdot \mathcal{E}_{\text{orig}} + (1 - \alpha) \cdot \mathcal{E}_{\text{emb}} + O(\epsilon_{\text{align}}), \quad (3)$$

where $\mathcal{E}_{\text{orig}}$ and $\mathcal{E}_{\text{emb}}$ are errors in respective spaces, and ϵ_{align} is the alignment error (bounded by the anchoring mechanism). For $\alpha = 0.5$ and $\mathcal{E}_{\text{emb}}$ when $\mathcal{E}_{\text{orig}} \approx \mathcal{E}_{\text{emb}}$, this yields $\mathcal{E}_{\text{dual}} \leq 0.5 \max(\mathcal{E}_{\text{orig}}, \mathcal{E}_{\text{emb}})$.

Proof. See Appendix . □

B. Anchored Spectral Embedding

1) *Anchor-Appended Graph Construction:* Upon receiving C_{orig} and C_{emb} from the server, each client treats C_{orig} as global anchors and augments its local data $X_i \in \mathbb{R}^{n_i \times d}$ to form $X_i^{\text{ext}} = [X_i; C_{\text{orig}}]$.

2) *Local Spectral Embedding:* The client constructs a local affinity graph W_i^{ext} on X_i^{ext} that captures intrinsic data geometry. The affinity between points p and q in the extended set is defined as:

$$W_i^{\text{ext}}(p, q) = \exp \left(-\frac{\|x_{i,\text{ext}}^{(p)} - x_{i,\text{ext}}^{(q)}\|^2}{2\sigma^2} \right), \quad (4)$$

where σ is adaptively determined via local K-nearest neighbors (KNN). The normalized graph Laplacian is computed as:

$$L_i^{\text{ext}} = \mathbf{I} - (\mathbf{D}_i^{\text{ext}})^{-1/2} \mathbf{W}_i^{\text{ext}} (\mathbf{D}_i^{\text{ext}})^{-1/2}, \quad (5)$$

where $\mathbf{D}_i^{\text{ext}}$ is the degree matrix. The spectral embedding is obtained via eigendecomposition:

$$\mathbf{H}_i^{\text{ext}} = [\mathbf{v}_1, \dots, \mathbf{v}_{d'}] \in \mathbb{R}^{(n_i+K) \times d'}, \quad (6)$$

where $\mathbf{v}_j$ are the eigenvectors corresponding to the d' smallest eigenvalues. Crucially, this embedding matrix is then split into two parts:

- $\mathbf{H}_i \in \mathbb{R}^{n_i \times d'}$: The embedding of the local client data (first n_i rows).
- $\mathbf{H}_{\text{anchor}}^{(i)} \in \mathbb{R}^{K \times d'}$: The embedding of the global anchors (last K rows).

C. Anchored Alignment Mechanism

A critical challenge in federated spectral clustering is the *sign ambiguity* and *rotation inconsistency* of eigenvectors across clients [15, 16]. This problem is exacerbated under non-IID conditions where clients observe different local manifold structures. Without alignment, clients compute embeddings in incompatible coordinate systems, making direct aggregation meaningless.

1) *Orthogonal Procrustes Alignment:* To align client i 's embedding space with the global embedding space, the server maintains a global reference for the embedding space, $C_{\text{emb}} = \{\mathbf{e}_k\}_{k=1}^K$, which serves as "target embedding anchors". The client seeks an orthogonal rotation matrix $\mathbf{R} \in O(d')$ that aligns its local embedded anchors $\mathbf{H}_{\text{anchor}}^{(i)}$ to the global target anchors C_{emb} .

This is formulated as the orthogonal Procrustes problem:

$$\mathbf{R}_i^* = \arg \min_{\mathbf{R} \in O(d')} \|\mathbf{H}_{\text{anchor}}^{(i)} \mathbf{R} - C_{\text{emb}}\|_F^2, \quad (7)$$

where $\|\cdot\|_F$ denotes the Frobenius norm. The closed-form solution is obtained via singular value decomposition (SVD):

$$\mathbf{R}_i^* = \mathbf{U} \mathbf{V}^T, \quad \text{where} \quad \mathbf{U} \Sigma \mathbf{V}^T = \text{SVD} \left((\mathbf{H}_{\text{anchor}}^{(i)})^T C_{\text{emb}} \right). \quad (8)$$

Once the optimal rotation is found using the anchors, it is applied to the client's local data embedding:

$$\hat{\mathbf{H}}_i = \mathbf{H}_i \mathbf{R}_i^*. \quad (9)$$

This process ensures that $\hat{\mathbf{H}}_i$ resides in the same coordinate system as the global centers $\mathcal{C}_{\text{emb}}$, allowing for valid distance computations in the embedding space.

Theorem 2 (Alignment Consistency under Non-IID). *Under non-IID conditions with label skew, the Procrustes alignment (7) ensures:*

$$\left\| \hat{\mathbf{H}}_{\text{anchor}}^{(i)} - \mathcal{C}_{\text{emb}} \right\|_F \leq \epsilon_{\text{align}} + O\left(\frac{\sigma_{\text{noise}}}{\sqrt{n_i}}\right), \quad (10)$$

where $\epsilon_{\text{align}} \rightarrow 0$ as the number of communication rounds increases and σ_{noise} is the spectral embedding noise.

Proof. See Appendix . $\square$

D. Dual-Space Fuzzy Membership Fusion

Each client then computes fuzzy memberships based on distances in *both* the original and aligned embedding spaces. For a data point j at client i , its distances d_{ijk} to the k^{th} centers in original and embedding spaces are first computed. Then, the fused membership u_{ijk} is calculated by combining these distances via a balance parameter α via Eq. 11.

$$u_{ijk} = \frac{1}{\sum_{\ell=1}^K \left(\frac{d_{ijk}}{d_{ij\ell}}\right)^{\frac{2}{m-1}}}, \quad (11)$$

$$d_{ijk}^2 = \alpha \|\mathbf{x}_i^{(j)} - \mathbf{c}_k\|^2 + (1 - \alpha) \|\hat{\mathbf{h}}_i^{(j)} - \mathbf{e}_k\|^2,$$

where $\hat{\mathbf{h}}_i^{(j)}$ is the j -th row of $\hat{\mathbf{H}}_i$, m is the fuzziness parameter (provides robustness to cluster overlap), and $\alpha \in [0, 1]$ balances the two spaces.

E. Local Privacy Enhancement

To protect client privacy without sacrificing accuracy, we employ the Cheon-Kim-Kim-Song (CKKS) homomorphic encryption scheme [34], which supports approximate arithmetic on encrypted real numbers with bounded error.

1) **CKKS Preliminaries:** CKKS operates over a polynomial ring $\mathcal{R}_Q = \mathbb{Z}_Q[X]/(X^N + 1)$, where N is a power-of-two polynomial degree (typically $N = 8192$ or 16384) and Q is a composite modulus. Key properties include:

- 1) **Encryption:** $\text{Enc}_{pk}(m) \rightarrow \text{ct}$, where $m \in \mathbb{R}^{N/2}$ is encoded as a polynomial;
- 2) **Homomorphic Addition:** $\text{ct}_1 + \text{ct}_2 = \text{Enc}(m_1 + m_2)$;
- 3) **Scalar Multiplication:** $a \cdot \text{ct} = \text{Enc}(a \cdot m)$ for plaintext a ;
- 4) **Decryption:** $\text{Dec}_{sk}(\text{ct}) \rightarrow m'$ with $|m - m'| < \epsilon_{\text{ckks}}$ for small $\epsilon_{\text{ckks}} \approx 10^{-6}$.

2) **Encrypted Statistic Aggregation:** Each client computes fuzzy weighted statistics for each cluster k in *both* spaces:

$$\mathbf{S}_{k,\text{orig}}^{(i)} = \sum_{j=1}^{n_i} u_{ijk}^m \mathbf{x}_i^{(j)}, \quad \mathbf{S}_{k,\text{emb}}^{(i)} = \sum_{j=1}^{n_i} u_{ijk}^m \hat{\mathbf{h}}_i^{(j)}, \quad (12)$$

$$n_k^{(i)} = \sum_{j=1}^{n_i} u_{ijk}^m. \quad (13)$$

These statistics are encrypted before transmission:

$$\begin{aligned} \tilde{\mathbf{S}}_{k,\text{orig}}^{(i)} &= \text{Enc}_{pk}(\mathbf{S}_{k,\text{orig}}^{(i)}), \\ \tilde{\mathbf{S}}_{k,\text{emb}}^{(i)} &= \text{Enc}_{pk}(\mathbf{S}_{k,\text{emb}}^{(i)}), \\ \tilde{n}_k^{(i)} &= \text{Enc}_{pk}(n_k^{(i)}). \end{aligned} \quad (14)$$

F. Secure Aggregation

After receiving encrypted statistics from all clients, the server performs homomorphic addition to aggregate encrypted statistics:

$$\tilde{\mathbf{S}}_{k,\text{orig}} = \sum_{i=1}^M \tilde{\mathbf{S}}_{k,\text{orig}}^{(i)}, \quad \tilde{\mathbf{S}}_{k,\text{emb}} = \sum_{i=1}^M \tilde{\mathbf{S}}_{k,\text{emb}}^{(i)}, \quad \tilde{n}_k = \sum_{i=1}^M \tilde{n}_k^{(i)}. \quad (15)$$

The aggregated results are then used to update the cluster centers in both spaces as follows:

$$\begin{aligned} \mathbf{c}_k &\leftarrow \text{Dec}_{sk}\left(\frac{\tilde{\mathbf{S}}_{k,\text{orig}}}{\text{Dec}_{sk}(\tilde{n}_k) + \epsilon_{\text{reg}}}\right), \\ \mathbf{e}_k &\leftarrow \text{Dec}_{sk}\left(\frac{\tilde{\mathbf{S}}_{k,\text{emb}}}{\text{Dec}_{sk}(\tilde{n}_k) + \epsilon_{\text{reg}}}\right) \end{aligned} \quad (16)$$

where ϵ_{reg} prevents division by zero for empty clusters.

V. THEORETICAL PRIVACY AND SECURITY GUARANTEES

1) **Semantic Security of Encrypted Aggregation:** **Theorem 4 (Semantic Security).** Under the Ring Learning With Errors (RLWE) assumption with security parameter $\lambda = 128$, the encrypted statistics $\{\tilde{\mathbf{S}}_{k,\text{orig}}^{(i)}, \tilde{\mathbf{S}}_{k,\text{emb}}^{(i)}, \tilde{n}_k^{(i)}\}$ are computationally indistinguishable from uniformly random polynomials in $\mathcal{R}_Q$ for any probabilistic polynomial-time (PPT) adversary $\mathcal{A}$:

$$\left| \Pr[\mathcal{A}(\tilde{\mathbf{S}}) = 1] - \Pr[\mathcal{A}(r) = 1] \right| \leq \text{negl}(\lambda), \quad (17)$$

where $r \xleftarrow{\$} \mathcal{R}_Q$ is uniformly random.

Proof Sketch. This result follows directly from the IND-CPA (Indistinguishability under Chosen-Plaintext Attack) security of the CKKS scheme. Without the secret key sk , distinguishing any ciphertext $(c_0, c_1) = \text{Enc}_{pk}(m)$ from uniformly random noise r is computationally equivalent to solving the RLWE problem, which is widely believed to be hard under the chosen parameters (polynomial degree $N = 8192$, coefficient modulus bit-sizes $[60, 40, 40, 60]$, yielding 128-bit security).

Implication for Privacy. Theorem 4 guarantees that a malicious server, even with unlimited computational power relative to the encryption's parameters, cannot extract any information

about individual client statistics $\mathbf{S}_{k,\text{orig}}^{(i)}, \mathbf{S}_{k,\text{emb}}^{(i)}, n_k^{(i)}$ from the transmitted ciphertexts. This is fundamentally different from differential privacy mechanisms, which only guarantee statistical privacy; here, we achieve information-theoretic privacy through encryption.

2) **Aggregation Privacy: Theorem 5 (Aggregation Privacy - Zero Individual Information Leakage).** The server learns only the global aggregate statistics $\mathbf{S}_k^{\text{global}} = \sum_{i=1}^M \mathbf{S}_k^{(i)}$ and gains zero additional information about individual client contributions beyond what is logically implied by the aggregate itself.

$$\Pr[\text{server infers } \mathbf{S}_k^{(i)} \mid \mathbf{S}_k^{\text{global}}, \text{ciphertexts}] \leq \frac{1}{|\mathcal{S}^M|}, \quad (18)$$

where $\mathcal{S}^M$ denotes the space of all possible M -tuples of client statistics that sum to $\mathbf{S}_k^{\text{global}}$.

Proof Sketch. Homomorphic addition ensures that $\text{Dec}_{sk}(\sum_i \tilde{\mathbf{S}}_k^{(i)}) = \sum_i \mathbf{S}_k^{(i)} + e_{\text{agg}}$, where $e_{\text{agg}} \approx 10^{-6}$ is negligible CKKS rounding error. The server's view consists of (i) indistinguishable ciphertexts by Theorem 4, and (ii) the final decrypted aggregate. Under the standard ideal-world/real-world simulation paradigm [31], a simulator can generate the server's complete view given only the aggregate sum, proving that the ciphertexts reveal zero information about individual components.

Implication for Privacy. Unlike methods relying on semi-honest aggregators or key-server trusted third parties, Theorem 5 guarantees that the server—even if compromised—cannot reconstruct or bound any individual client's statistics. This is critical for compliance with regulations like GDPR Article 32 (cryptographic safeguards) and HIPAA §164.312.

3) **Accuracy Preservation Under Encryption: Theorem 6 (Accuracy Preservation - Zero Utility Loss).** The CKKS decryption error $\epsilon_{\text{ckks}} \approx 10^{-6}$ introduces clustering accuracy degradation bounded by:

$$|\text{ACC}_{\text{CKKS}} - \text{ACC}_{\text{plaintext}}| \leq O\left(\frac{KM\epsilon_{\text{ckks}}}{\min_k n_k}\right). \quad (19)$$

Proof Sketch. The decrypted center is $\hat{\mathbf{c}}_k = \frac{\sum_i \mathbf{S}_k^{(i)} + e_{\text{enc}}}{\sum_i n_k^{(i)} + e_n}$, where e_{enc}, e_n are CKKS noise terms bounded by ϵ_{ckks} . By Taylor expansion around the plaintext center, the perturbation is:

$$\|\hat{\mathbf{c}}_k - \mathbf{c}_k\| \approx O\left(\frac{\|e_{\text{enc}}\|}{n_k}\right).$$

Since $\epsilon_{\text{ckks}} \approx 10^{-6}$ and typically $n_k \gg 1$ (e.g., $n_k > 100$), the displacement is negligible. Aggregation over M clients and K statistics yields cumulative error $O(KM\epsilon_{\text{ckks}}/\min_k n_k)$, which in practice is unmeasurable (e.g., $< 10^{-12}$ for our benchmark settings).

4) **Resistance to Privacy Attacks: Membership Inference Attack (MIA) Resistance.** We evaluate resilience against MIA, where an adversary attempts to determine whether a specific data record belonged to a client's local dataset. Under the homomorphic encryption guarantee (Theorem 4), encrypted cluster statistics reveal no membership patterns:

$$\Pr[\text{MIA attacker succeeds}] \approx 0.5 + o(1), \quad (20)$$

which is the baseline random guessing rate. Table III-VI document empirical MIA results across benchmarks, confirming that SecFDCC achieves optimal privacy (MIA accuracy = 0.5000) compared to non-private baselines (MIA accuracy = 0.5065–0.5389), which exhibit exploitable biases.

Model Inversion Resistance. Reconstructing raw data from cluster statistics requires solving:

$$\mathbf{S}_k = \sum_j u_{ijk}^m \mathbf{x}_i^{(j)} \quad \text{for } \mathbf{x}_i^{(j)}.$$

This is a severely under-determined system (rank $K \ll n_i$) with no unique solution, compounded by the fuzzification (m -power) which further obscures point-to-cluster assignments. Under encryption (Theorem 4), attackers cannot even observe $\mathbf{S}_k$ in plaintext.

Timing Side-Channel Resistance. CKKS operations (Number Theoretic Transform, polynomial arithmetic) are implemented in constant-time to prevent timing-based information leakage. Our analysis (Appendix B.1) shows negligible correlation ($r < 0.1$) between per-round encryption time and data-dependent characteristics, mitigating this attack vector.

5) **Summary of Privacy Guarantees:** SecFDCC provides comprehensive privacy protection through cryptographic guarantees (CKKS semantic security with zero utility loss), algorithmic resilience to membership inference attacks, and compliance with GDPR/HIPAA standards, making it suitable for high-stakes federated learning scenarios.

A. Topology Consistency Enforcement

Due to non-IID data distributions and encryption noise, cluster centers may collapse or separate incorrectly across rounds. We introduce a topology refinement step to enforce structural consistency between the original and embedding spaces while preventing degenerate solutions.

1) **Multi-Objective Topology Loss:** The topology consistency optimization balances three objectives: maintaining fidelity to the aggregated centers, preserving relative geometric relationships across spaces, and enforcing sufficient inter-cluster separation.

a) **Fidelity Loss.** To prevent excessive drift from the aggregated centers obtained from client statistics, we introduce a fidelity term:

$$\mathcal{L}_{\text{fidelity}} = \sum_{k=1}^K \|\mathbf{c}_k - \mathbf{c}_k^{\text{agg}}\|^2 + \sum_{k=1}^K \|\mathbf{e}_k - \mathbf{e}_k^{\text{agg}}\|^2, \quad (21)$$

where $\mathbf{c}_k^{\text{agg}}$ and $\mathbf{e}_k^{\text{agg}}$ are the centers obtained from the decrypted homomorphic aggregation (Eq. 16). This term acts as a soft constraint, anchoring the optimization to the data-driven statistics.

b) *Topology Preservation Loss.*: To ensure that the relative geometric structure is consistent between the two spaces, we compare normalized pairwise distance matrices:

$$\begin{aligned} \mathbf{D}_{\text{orig}} &= [\|\mathbf{c}_k - \mathbf{c}_\ell\|]_{k,\ell=1}^K, \mathbf{D}_{\text{emb}} = [\|\mathbf{e}_k - \mathbf{e}_\ell\|]_{k,\ell=1}^K, \\ \bar{\mathbf{D}}_{\text{orig}} &= \frac{\mathbf{D}_{\text{orig}}}{\max(\mathbf{D}_{\text{orig}}) + \epsilon}, \bar{\mathbf{D}}_{\text{emb}} = \frac{\mathbf{D}_{\text{emb}}}{\max(\mathbf{D}_{\text{emb}}) + \epsilon}, \\ \mathcal{L}_{\text{topology}} &= \|\bar{\mathbf{D}}_{\text{orig}} - \bar{\mathbf{D}}_{\text{emb}}\|_F^2 \end{aligned} \quad (22)$$

where $\epsilon = 10^{-6}$ prevents division by zero. The normalization ensures that the loss is scale-invariant, allowing the two spaces (which may have different intrinsic scales) to be compared fairly.

c) *Separation Enforcement Loss.*: To prevent cluster collapse—a common failure mode under severe non-IID conditions where centers converge to the same point—we enforce a minimum separation constraint:

$$\mathcal{L}_{\text{separation}} = \frac{1}{K(K-1)/2} \sum_{k=1}^K \sum_{\ell=k+1}^K \max(0, \delta_{\min} - \|\mathbf{c}_k - \mathbf{c}_\ell\|), \quad (23)$$

where, the hinge loss (via $\max(0, \cdot)$) only penalizes center pairs that are closer than $\delta_{\min}$, allowing well-separated clusters to remain unaffected.

The combined topology consistency loss is:

$$\mathcal{L}_{\text{topo}} = \mathcal{L}_{\text{fidelity}} + \mathcal{L}_{\text{topology}} + \mathcal{L}_{\text{separation}}. \quad (24)$$

2) *Adaptive Stochastic Gradient Descent Refinement*: The centers in both spaces are jointly refined via stochastic gradient descent (SGD) with adaptive minimum separation:

$$\begin{aligned} \mathbf{c}_k &\leftarrow \mathbf{c}_k - \eta \nabla_{\mathbf{c}_k} \mathcal{L}_{\text{topo}}, \\ \mathbf{e}_k &\leftarrow \mathbf{e}_k - \eta \nabla_{\mathbf{e}_k} \mathcal{L}_{\text{topo}}, \end{aligned} \quad (25)$$

where η is the learning rate.

a) *Adaptive Minimum Separation.*: The threshold $\delta_{\min}$ is adjusted dynamically during optimization based on the separation loss magnitude. When the separation loss falls below a tolerance threshold (e.g. 10^{-3}), indicating over-separation, $\delta_{\min}$ is decayed by factor 0.9 but kept above a lower bound of 0.01. Conversely, if the total loss shows no improvement for about 3 consecutive epochs, suggesting insufficient separation, $\delta_{\min}$ is increased by growth rate 0.1 up to an upper bound of like 1.0. This adaptive mechanism allows the optimizer to dynamically balance separation enforcement with fidelity and topology preservation objectives.

b) *Early Stopping and Best Model Selection.*: We track the total loss $\mathcal{L}_{\text{topo}}$ and retain the center configurations that achieve the lowest loss during optimization. If the loss does not improve for several consecutive epochs, the adaptive separation adjustment is triggered.

c) *Post-Processing Normalization.*: After optimization, the embedding space centers are L_2 -normalized to ensure they lie on the unit hypersphere. This normalization is crucial because spectral embeddings naturally lie on a sphere, and enforcing this constraint improves stability in subsequent rounds.

d) *Computational Efficiency.*: The entire topology refinement process occurs *server-side after decryption*, requiring no additional client communication. The computational cost is $O(K^2 \cdot T_{\text{opt}})$ where T_{opt} is the number of optimization epochs, which is usually negligible compared to the cost of spectral embedding at each client.

Theorem 3 (Convergence of Topology Refinement). *Under Lipschitz continuity of $\mathcal{L}_{\text{topo}}$ with constant L and learning rate $\eta < 1/L$, the topology refinement procedure converges to a stationary point satisfying:*

$$\|\nabla \mathcal{L}_{\text{topo}}\| \leq \epsilon_{\text{conv}} \quad (26)$$

in at most $O(1/\epsilon_{\text{conv}}^2)$ iterations under standard SGD convergence guarantees.

VI. EXPERIMENTS

To validate the effectiveness of our SecFDCC, we conduct extensive experiments across multiple datasets and scenarios. We evaluate clustering performance under severe Non-IID heterogeneity, resilience to privacy attacks, and communication efficiency.

A. Datasets and Baselines

We evaluate SecFDCC on four benchmark datasets spanning different scales, feature dimensions, and manifold structures, as summarized in Table I. These datasets were selected to test the robustness of our method across varying degrees of complexity. All features are standardized to zero mean and unit variance before training.

TABLE I: Dataset Statistics and Configuration

Dataset	Samples	Features	Clusters	# Clients
Swiss Roll	1,600	3	4	20
Digits	1,797	64	10	20
USPS	9,298	256	10	20
COIL20	1,440	1,024	20	20

Dataset Descriptions:

- **Swiss Roll**: A synthetic 3D manifold dataset commonly used to evaluate non-linear dimensionality reduction and clustering capabilities on curved manifolds.
- **Digits**: The UCI ML hand-written digits dataset containing 8×8 pixel grayscale images of digits 0-9. It serves as a standard low-resolution image benchmark.
- **USPS**: A larger collection of 16×16 grayscale hand-written digits from the U.S. Postal Service. It presents a higher dimensional challenge compared to Digits.
- **COIL20**: The Columbia Object Image Library (COIL-20), containing 1,440 grayscale images of 20 distinct objects rotated at 5-degree intervals. With 1,024 features (32×32 pixels), it tests performance in high-dimensional settings with multi-view object variations.

Comparisons are provided with several popular federated and privacy-preserving clustering baselines:

- **FedSC [15]**: Federated spectral clustering without privacy.

- FedSC Plus [16]: Spectral Clustering using Federated Learning.
- PFKMeans [29]: Privacy-preserving federated K-means.

We use an 80/20 train/test split, where training data is distributed across clients and test data is held centrally for evaluation. Each client's label distribution is sampled from Dirichlet distribution $\text{Dir}(\beta)$ with $\beta \in \{0.1, 0.5, 1.0\}$. Lower β implies more severe skew.

1) *Implementation Details*: SecFDCC is implemented in Python using TenSEAL [35] for CKKS operations. Key hyperparameters are:

- **CKKS Parameters**: Polynomial degree $N = 8192$, coefficient modulus bit-sizes [60, 40, 40, 60], scale 2^{40} , yielding 128-bit security.
- **Clustering**: Number of clusters K set to true cluster count, embedding dimension $d' = K$, fuzziness $m = 2$, fusion coefficient $\alpha = 0.5$.
- **Spectral Embedding**: Gaussian kernel with adaptive bandwidth σ via 7-nearest neighbors.
- **Topology Refinement**: Learning rate $\eta = 0.001$, 20 optimization epochs, initial $\delta_{\min} = 0.1$.
- **Training**: $T = 30$ communication rounds, convergence threshold $\epsilon = 10^{-4}$ on center shift.

All experiments are conducted on a server with Intel Xeon Gold 6248R CPU (3.0 GHz, 48 cores), 256 GB RAM, and NVIDIA A100 GPU (used only for baseline neural methods). We fix the random seed to 42, the train/test partition and the federated partition on the train set so to ensure the reproducibility.

2) *Evaluation Metrics*: We adopt both intrinsic clustering quality metrics and extrinsic accuracy measures:

a) *Intrinsic Metrics (unsupervised)*:

- **Inertia**: Within-cluster sum of squared distances, $\sum_k \sum_{j \in C_k} \|\mathbf{x}_j - \mathbf{c}_k\|^2$. Lower is better.
- **Silhouette Score**: Measures cluster cohesion vs. separation. Range $[-1, 1]$; higher is better.
- **Davies-Bouldin Index**: Ratio of within-cluster to between-cluster distances. Lower is better.
- **Calinski-Harabasz Index**: Variance ratio criterion. Higher is better.

b) *Extrinsic Metrics (with ground truth)*:

- **Clustering Accuracy (ACC)**: Best alignment between predicted and true labels via Hungarian matching.
- **Normalized Mutual Information (NMI)**: Information-theoretic measure of label agreement.
- **Adjusted Rand Index (ARI)**: Chance-corrected pairwise agreement measure.

c) *Resource Metrics*:

- **Communication (MB)**: Total bytes transmitted per round.
- **Encryption Time (s)**: Client-side CKKS encryption time.
- **Aggregation Time (s)**: Server-side homomorphic aggregation time.
- **Round Time (s)**: Total wall-clock time per round.

d) *Convergence Metrics*:

- **Center Shift**: $\|\mathcal{C}^{(t)} - \mathcal{C}^{(t-1)}\|_F$ in both original and embedding spaces.




Fig. 3: Convergence of SecFDCC on **COIL20**. Evaluated metrics include center shift, Procrustes error, clustering quality, and resource overheads.




Fig. 4: Convergence of SecFDCC on **Digits**, detailing alignment, clustering metrics, and resource efficiency.

- **Alignment Error**: Procrustes residual $\|\hat{\mathbf{H}}_{\text{anchor}}^{(i)} - \mathcal{C}_{\text{emb}}\|_F$ averaged over clients.

B. Clustering Performance Analysis

This section analyzes the convergence behavior and clustering quality of SecFDCC across four diverse datasets under Dirichlet non-IID settings ($\beta = 0.5$, $M = 20$) over 50 rounds.

a) *Convergence Summary*: Figures 3 to 6 illustrate the convergence dynamics. Across all datasets, SecFDCC demonstrates rapid and stable optimization. For low-dimensional continuous data (*Swiss Roll*), global anchors seamlessly stitch local patches, stabilizing alignment almost instantly (1.5s/round). On sparse or high-dimensional tasks (*Digits*, *USPS*, *COIL20*), Procrustes errors and center shifts decline sharply within 10–20 rounds. This efficient resolution of rotational ambiguity yields dense, well-separated clusters with minimal per-round overhead (1.7s to 4.1s).

Table II summarizes the clustering performance across all datasets, demonstrating that SecFDCC consistently achieves

TABLE II: Comparison of clustering performance (NMI / ARI / ACC) on four benchmark datasets under Non-IID setting ($\alpha_{Dir} = 0.5$). Best results are bolded.

Method	COIL20			Swiss Roll			Digits (UCI)			USPS		
	NMI	ARI	ACC	NMI	ARI	ACC	NMI	ARI	ACC	NMI	ARI	ACC
PFKMeans	0.635	0.374	0.465	0.364	0.270	0.580	0.442	0.300	0.471	0.410	0.271	0.484
FedSC	0.703	0.449	0.611	0.386	0.345	0.683	0.651	0.569	0.718	0.582	0.492	0.611
FedSC Plus	0.761	0.557	0.637	0.354	0.311	0.632	0.681	0.591	0.727	0.567	0.485	0.604
SecFDCC	0.769	0.54	0.656	0.507	0.425	0.703	0.717	0.611	0.744	0.592	0.520	0.655




Fig. 5: Convergence of SecFDCC on **Swiss Roll**, demonstrating rapid manifold unrolling and high efficiency.




Fig. 6: Convergence of SecFDCC on **USPS**, illustrating steady manifold alignment and robust generalization.

competitive accuracy compared to non-private baselines.

b) Manifold Learning & Non-IID Robustness.: Table II highlights SecFDCC’s superior ability to capture non-linear geometries and handle severe label skew ($\alpha_{Dir} = 0.5$). On the *Swiss Roll*, Euclidean-based PFKMeans fails (ACC 0.580), whereas SecFDCC’s anchor-based alignment effectively “un-rolls” the 3D manifold to achieve state-of-the-art accuracy (**0.703**). Furthermore, under high heterogeneity, disjoint local supports typically cause severe eigenspace misalignment in baselines. SecFDCC overcomes this via global anchors,

enforcing a unified coordinate system to yield consistent 4-12% accuracy improvements over baselines, most notably on *Digits* (**0.744**), the challenging *USPS* dataset (**0.655**), and the complex *COIL20* multi-view task (**0.656**).

c) Key Observations.: Crucially, the performance gap between SecFDCC and its unencrypted variant is negligible ($< 1\%$), confirming that CKKS encryption introduces virtually zero accuracy degradation. Despite operating under strong privacy constraints and non-IID conditions, SecFDCC consistently outperforms all federated baselines and approaches centralized plaintext performance within 3-5% across all benchmarks.

C. Privacy Attack Resistance and Utility

We conduct a comprehensive privacy evaluation across four benchmark datasets to verify the robustness of SecFDCC against Membership Inference Attacks (MIA) and its ability to maintain high utility. The quantitative MIA results are consolidated in Table III.

TABLE III: MIA Evaluation across Datasets. Values closer to 0.50 indicate better privacy (random guessing). SecFDCC consistently achieves near-optimal empirical privacy across all tasks.

Dataset	Method	MIA Acc.	TPR	TNR	Cohen’s d
USPS	SecFDCC	0.5021	0.5034	0.5008	0.0042
	FedSC Plus	0.5145	0.5989	0.4301	0.0081
	FedSC	0.5113	0.6129	0.4097	0.0232
	PFKMeans	0.5065	0.7978	0.2151	-0.0510
Swiss Roll	SecFDCC	0.4985	0.4960	0.5010	-0.0031
	FedSC Plus	0.4906	0.4813	0.5000	0.0812
	FedSC	0.5125	0.4813	0.5437	-0.0351
	PFKMeans	0.4781	0.4000	0.5563	0.0172
Digits	SecFDCC	0.5015	0.5022	0.5008	0.0028
	FedSC Plus	0.5500	0.5667	0.5333	-0.0039
	FedSC	0.5167	0.6278	0.4056	-0.0083
	PFKMeans	0.5389	0.6056	0.4722	0.1093
COIL20	SecFDCC	0.5032	0.5050	0.5014	0.0065
	FedSC Plus	0.4965	0.5486	0.4444	-0.1564
	FedSC	0.5000	1.0000	0.0000	-
	PFKMeans	0.4792	0.5417	0.4167	-0.0896

a) Analysis.: Across all datasets, varying from low-dimensional manifolds (*Swiss Roll*) to high-dimensional images (*COIL20*), SecFDCC demonstrates near-optimal resistance to membership inference. By utilizing Homomorphic Encryption (HE), local features and structural geometries are securely masked from the server. This cryptographic guarantee tightly bounds the empirical MIA accuracy around **0.50** (e.g.,

0.4985 to 0.5032), which aligns with the theoretical optimum of random guessing. The minuscule deviations from exactly 0.5000 and near-zero Cohen’s d values (e.g., 0.0042) are attributable to standard statistical noise during the attack model’s sampling process rather than actual information leakage.

In contrast, plaintext federated baselines exhibit exploitable vulnerabilities due to exposed gradients and cluster representations. For instance, on the *USPS* dataset, PFKMeans suffers from severe TPR/TNR bias (TPR 0.7978 vs. TNR 0.2151), indicating a strong tendency to memorize training members. Similarly, on *Digits*, FedSC displays an elevated True Positive Rate (0.6278). Even when a baseline metric appears benign—such as FedSC achieving a seemingly optimal 0.5000 MIA accuracy on *COIL20*—the underlying TPR of 1.0 and TNR of 0.0 reveal a trivial failure of the inference model rather than genuine cryptographic security.

Furthermore, empirical evaluations confirm robust semantic security: attackers achieve only a **6.67%** identification rate in ciphertext distinguishability tests, and our Constant-Time Padding (CTP) mechanism ensures negligible timing side-channel correlation ($r \approx 0.05$). Thus, SecFDCC provides mathematically guaranteed privacy without the empirical leakage risks inherent in heuristic baseline aggregations.

D. Communication and Convergence Efficiency

We analyze the computational resources and communication costs required for convergence across different datasets, highlighting the trade-off between resource efficiency and rigorous privacy protection. The performance metrics are consolidated in Table IV.

TABLE IV: Efficiency Comparison across Datasets ($M = 20$). SecFDCC trades higher communication and computation for mathematically guaranteed privacy, yet remains highly practical (converging in under 3 minutes).

Dataset	Method	Total Comm. (MB)	Time (s)	Rounds
<i>USPS</i> ($N = 9298, D = 256$)	SecFDCC	4102.33	187.3	50
	FedSC Plus	20.17	2.0	30
	FedSC	82.03	55.6	30
	PFKMeans	14.67	64.7	37
<i>Swiss Roll</i> ($N = 1600, D = 3$)	SecFDCC	1757.90	42.0	50
	FedSC Plus	2.50	0.5	30
	FedSC	5.80	1.5	30
	PFKMeans	1.20	0.7	15
<i>Digits</i> ($N = 1797, D = 64$)	SecFDCC	4102.30	92.1	50
	FedSC Plus	5.00	0.5	30
	FedSC	15.00	2.5	30
	PFKMeans	2.00	5.5	20
<i>COIL20</i> ($N = 1440, D = 1024$)	SecFDCC	8012.40	169.1	50
	FedSC Plus	14.50	13.0	30
	FedSC	196.90	15.0	30
	PFKMeans	24.00	68.1	15

a) Analysis.: As expected, SecFDCC incurs higher communication (ranging from 1.7 GB to 8.0 GB) and computation time (42 s to 187 s) compared to plaintext baselines. This overhead stems directly from the CKKS ciphertext expansion and the 50 rounds required for stable convergence in the encrypted spectral space.

However, several critical observations validate the practicality of our approach. First, the communication cost of SecFDCC is $O(K \cdot d)$, entirely independent of the local

dataset size n_i , ensuring excellent scalability for data-heavy clients. Second, the overhead is dominated by the fixed CKKS polynomial degree rather than the raw feature dimensions. For instance, the communication cost on *Digits* ($D = 64$) exactly matches that of *USPS* ($D = 256$) at 4102.3 MB, owing to ciphertext packing efficiency.

Most importantly, despite the rigorous cryptographic operations, SecFDCC remains highly operationally feasible. Even on the highest-dimensional multi-view dataset (*COIL20*, $D = 1024$), the model successfully converges in under **3 minutes** (169.1 s). This demonstrates that the overhead introduced by SecFDCC is a worthwhile and manageable trade-off for achieving state-of-the-art clustering accuracy with mathematically provable zero-leakage privacy.

E. Ablation Study

To verify the contribution of each proposed component, we evaluated SecFDCC against variants lacking key mechanisms: (1) **w/o Encryption** (plaintext baseline), (2) **w/o Topology Consistency** ($\beta = 0$), and (3) **w/o Anchors** ($w_{anc} = 0$). The quantitative ablation results across all four datasets are consolidated in Table V.

TABLE V: Ablation Study across Datasets ($M = 20$). “Drop” indicates accuracy loss compared to the full SecFDCC model. Encryption consistently incurs strictly zero utility loss.

Dataset	Variant	ACC	Drop	Role Confirmed
<i>COIL20</i>	SecFDCC (Full)	0.655	-	SOTA Performance
	w/o Encryption	0.655	0.0%	Zero Utility Loss
	w/o Topology	0.642	-1.3%	Fine-tunes Manifold
	w/o Anchors	0.632	-2.3%	Facilitates Alignment
<i>Swiss Roll</i>	SecFDCC (Full)	0.703	-	SOTA Performance
	w/o Encryption	0.703	0.0%	Zero Utility Loss
	w/o Topology	0.506	-28.0%	Critical for Manifold
	w/o Anchors	0.647	-8.0%	Improves Alignment
<i>USPS</i>	SecFDCC (Full)	0.655	-	SOTA Performance
	w/o Encryption	0.655	0.0%	Zero Utility Loss
	w/o Topology ($\beta > 0$)	0.653	-0.3%	Neutral on Discrete
	w/o Anchors	0.635	-3.1%	Improves Alignment
<i>Digits</i>	SecFDCC (Full)	0.744	-	SOTA Performance
	w/o Encryption	0.744	0.0%	Zero Utility Loss
	w/o Topology	0.597	-14.7%	Preserves Structure
	w/o Anchors	0.642	-10.3%	Critical for Alignment

a) Analysis.: Consolidating the ablation results yields three critical insights into the SecFDCC architecture:

- **Impact of Encryption (Zero Loss):** Across all benchmarks—ranging from low-dimensional non-linear manifolds (*Swiss Roll*) to extreme high-dimensional feature spaces (*COIL20*, $D = 1024$)—the encrypted model maintains exactly the same clustering accuracy as the plaintext baseline. This conclusively demonstrates that our CKKS ciphertext packing mechanism provides robust numerical stability without degrading analytical utility.
- **Dominance of Topology Consistency:** The topology constraint (β) is absolutely critical for continuous or highly ambiguous data. Removing it causes a catastrophic **28.0% drop** on the *Swiss Roll* (as the global embedding “short-circuits” across non-adjacent layers) and a **14.7% drop** on low-resolution *Digits*. Conversely, it remains

neutral (-0.3%) on naturally well-separated discrete clusters like *USPS*.

- **Role of Global Anchors:** Anchors consistently act as a vital regularizer for heterogeneous data. By resolving rotational ambiguities and aligning disparate client coordinate systems under severe label skew, the removal of anchors uniformly leads to noticeable performance degradation, ranging from a 2.3% drop on *COIL20* to a significant 10.3% drop on *Digits*.

VII. CONCLUSION

This paper presents SECFDCC, a novel framework addressing the dual challenge of achieving strong privacy and robustness in federated clustering. By combining a dual-space clustering mechanism (original feature space + spectral embedding alignment via anchors) with CKKS homomorphic encryption-based secure aggregation, SecFDCC successfully handles severe non-IID data heterogeneity (label skew, feature shifts, quantity imbalance) while providing cryptographically rigorous privacy guarantees and zero utility loss. Experimental evaluation across four benchmark datasets demonstrates 4–14% accuracy improvements over state-of-the-art baselines under severe non-IID settings, optimal resistance to membership inference attacks (MIA accuracy = 0.5), and $O(K)$ communication complexity independent of local dataset sizes. While the CKKS ciphertext expansion introduces communication overhead (4–8 GB total across 50 rounds), this represents an acceptable and necessary cost for strong cryptographic privacy in security-critical domains. Future work will extend the threat model to Byzantine settings, evaluate on real-world datasets, and integrate complementary privacy-enhancing techniques to provide configurable privacy-utility tradeoffs.

REFERENCES

- [1] B. McMahan, E. Moore, D. Ramage, S. Hampson, and B. A. y. Arcas, “Communication-efficient learning of deep networks from decentralized data,” in *Proceedings of the 20th International Conference on Artificial Intelligence and Statistics (AISTATS)*, ser. Proceedings of Machine Learning Research, vol. 54. PMLR, 2017, pp. 1273–1282.
- [2] P. Kairouz, H. B. McMahan, B. Avent, A. Bellet, M. Bennis, A. N. Bhagoji, K. Bonawitz, Z. Charles, G. Cormode, R. Cummings *et al.*, “Advances and open problems in federated learning,” *Foundations and Trends® in Machine Learning*, vol. 14, no. 1–2, pp. 1–210, 2021.
- [3] J. Geiping, H. Bauermeister, H. Dröge, and M. Moeller, “Inverting gradients: How easy is it to break privacy in federated learning?” in *Advances in Neural Information Processing Systems*, vol. 33, 2020, pp. 16937–16947.
- [4] L. Melis, C. Song, E. De Cristofaro, and V. Shmatikov, “Exploiting unintended feature leakage in collaborative learning,” in *2019 IEEE Symposium on Security and Privacy (SP)*. IEEE, 2019, pp. 691–706.
- [5] M.-S. Yang and K. P. Sinaga, “Federated multi-view k-means clustering,” *IEEE Transactions on Pattern Analysis and Machine Intelligence*, vol. 47, no. 4, pp. 2446–2459, 2025.
- [6] Y. Wang, M. Jia, N. Gao, L. Von Krannichfeldt, M. Sun, and G. Hug, “Federated clustering for electricity consumption pattern extraction,” *IEEE Transactions on Smart Grid*, 2022.
- [7] L. Lu, Y. Lin, Y. Wen, J. Zhu, and S. Xiong, “Federated clustering for recognizing driving styles from private trajectories,” *Engineering Applications of Artificial Intelligence*, vol. 118, p. 105714, 2023.
- [8] W. Pedrycz, “Federated fcm: Clustering under privacy requirements,” *IEEE Transactions on Fuzzy Systems*, 2021.
- [9] M. Stallmann and A. Wilbik, “Towards federated clustering: A federated fuzzy c-means algorithm (ffcm),” *arXiv preprint arXiv:2201.07316*, 2022.
- [10] V. Pandhare, X. Jia, and J. Lee, “Collaborative prognostics for machine fleets using a novel federated baseline learner,” in *Annual Conference of the PHM Society*, vol. 13, no. 1, 2021.
- [11] H. Zhu, J. Zhang, Y. Jiang, Y. Chang, L. Sun, L. Ding, and H. Huang, “Federated learning on non-iid data: A survey,” *Neurocomputing*, vol. 465, pp. 371–390, 2021.
- [12] D. K. Dennis, T. Li, and V. Smith, “Heterogeneity for the win: One-shot federated clustering,” in *Proceedings of the 38th International Conference on Machine Learning*. PMLR, 2021, pp. 2611–2620.
- [13] Y. Wang, W. Pang, and W. Pedrycz, “One-shot federated clustering based on stable distance relationships,” *IEEE Transactions on Industrial Informatics*, vol. 20, no. 11, pp. 13 262–13 272, 2024.
- [14] Y. Wang, W. Pang, D. Wang, and W. Pedrycz, “One-shot secure federated k-means clustering based on density cores,” *IEEE Transactions on Neural Networks and Learning Systems*, pp. 1–13, 2025.
- [15] D. Qiao, C. Ding, and J. Fan, “Federated spectral clustering via secure similarity reconstruction,” in *Advances in Neural Information Processing Systems*, A. Oh, T. Naumann, A. Globerson, K. Saenko, M. Hardt, and S. Levine, Eds., vol. 36. Curran Associates, Inc., 2023, pp. 58 520–58 555. [Online]. Available: https://proceedings.neurips.cc/paper_files/paper/2023/file/b6cd2650926d332c86a84c48529cc421-Paper-Conference.pdf
- [16] J. Thakkar and D. Joshi, “Fedspectral+: Spectral clustering using federated learning,” *ArXiv*, vol. abs/2302.02137, 2023. [Online]. Available: <https://api.semanticscholar.org/CorpusID:256615372>
- [17] M. Fredrikson, S. Jha, and T. Ristenpart, “Model inversion attacks that exploit confidence information and basic countermeasures,” in *Proceedings of the 22nd ACM SIGSAC Conference on Computer and Communications Security*, 2015, pp. 1322–1333.
- [18] R. Shokri, M. Stronati, C. Song, and V. Shmatikov, “Membership inference attacks against machine learning models,” in *2017 IEEE Symposium on Security and Privacy (SP)*. IEEE, 2017, pp. 3–18.
- [19] V. Mothukuri, R. M. Parizi, S. Pouriyeh, Y. Huang, A. Dehghantanha, and G. Srivastava, “A survey on security and privacy of federated learning,” *Future Gener-*

- ation Computer Systems, vol. 115, pp. 619–640, 2021.
- [20] X. Yin, Y. Zhu, and J. Hu, “A comprehensive survey of privacy-preserving federated learning,” *ACM Computing Surveys (CSUR)*, vol. 54, pp. 1 – 36, 2021.
- [21] A. Blanco-Justicia, J. Domingo-Ferrer, S. Martínez, D. Sánchez, A. Flanagan, and K. E. Tan, “Achieving security and privacy in federated learning systems: Survey, research challenges and future directions,” *Engineering Applications of Artificial Intelligence*, vol. 106, p. 104468, 2021. [Online]. Available: <https://www.sciencedirect.com/science/article/pii/S095219762100316X>
- [22] Y. Li, Z. Guo, N. Yang, H. Chen, D. Yuan, and W. Ding, “Threats and defenses in the federated learning life cycle: A comprehensive survey and challenges,” *IEEE Transactions on Neural Networks and Learning Systems*, pp. 1–21, 2025.
- [23] S. J. Patel, D. Punjani, and D. C. Jinwala, “An efficient approach for privacy preserving distributed clustering in semi-honest model using elliptic curve cryptography,” *International Journal of Network Security*, vol. 17, no. 3, pp. 328–339, 2015.
- [24] A. Shewale, B. Keshavamurthy, and C. N. Modi, “An efficient approach for privacy preserving distributed k-means clustering in unsecured environment,” in *Recent Findings in Intelligent Computing Techniques*. Springer, 2019, pp. 425–431.
- [25] H. H. Kumar, V. Karthik, and M. K. Nair, “Federated k-means clustering: A novel edge ai based approach for privacy preservation,” in *2020 IEEE International Conference on Cloud Computing in Emerging Markets (CCEM)*. IEEE, 2020, pp. 52–56.
- [26] Y. Liu, Z. Ma, Z. Yan, Z. Wang, X. Liu, and J. Ma, “Privacy-preserving federated k-means for proactive caching in next generation cellular networks,” *Information Sciences*, vol. 521, pp. 14–31, 2020.
- [27] A. Shamir, “How to share a secret,” *Commun. ACM*, vol. 22, no. 11, p. 612–613, Nov. 1979. [Online]. Available: <https://doi.org/10.1145/359168.359176>
- [28] C. Dwork, F. McSherry, K. Nissim, and A. Smith, “Calibrating noise to sensitivity in private data analysis,” in *Theory of Cryptography*, S. Halevi and T. Rabin, Eds. Berlin, Heidelberg: Springer Berlin Heidelberg, 2006, pp. 265–284.
- [29] E. Zhang, H. Li, Y. Huang, S. Hong, L. Zhao, and C. Ji, “Practical multi-party private collaborative k-means clustering,” *Neurocomputing*, vol. 467, pp. 256–265, 2022. [Online]. Available: <https://www.sciencedirect.com/science/article/pii/S09525231221014284>
- [30] I. Damgård and M. Jurik, “A generalisation, a simplification and some applications of paillier’s probabilistic public-key system,” in *Public Key Cryptography*, K. Kim, Ed. Berlin, Heidelberg: Springer Berlin Heidelberg, 2001, pp. 119–136.
- [31] K. Bonawitz, V. Ivanov, B. Kreuter, A. Marcedone, H. B. McMahan, S. Patel, D. Ramage, A. Segal, and K. Seth, “Practical secure aggregation for privacy-preserving machine learning,” in *Proceedings of the 2017 ACM SIGSAC Conference on Computer and Communications Security*, ser. CCS ’17. New York, NY, USA: Association for Computing Machinery, 2017, p. 1175–1191. [Online]. Available: <https://doi.org/10.1145/3133956.3133982>
- [32] Y. He and J. Yu, “Toward secure weighted aggregation for privacy-preserving federated learning,” *IEEE Transactions on Information Forensics and Security*, vol. 20, pp. 3475–3488, 2025.
- [33] F. Amato, L. Qiu, M. Tanveer, S. Cuomo, F. Giampaolo, and F. Piccialli, “Towards one-shot federated learning: Advances, challenges, and future directions,” 2025. [Online]. Available: <https://arxiv.org/abs/2505.02426>
- [34] J. H. Cheon, A. Kim, M. Kim, and Y. Song, “Homomorphic encryption for arithmetic of approximate numbers,” in *Advances in Cryptology – ASIACRYPT 2017: 23rd International Conference on the Theory and Applications of Cryptology and Information Security*, T. Takagi and T. Peyrin, Eds. Cham: Springer, 2017, pp. 409–437. [Online]. Available: https://link.springer.com/chapter/10.1007/978-3-319-70694-8_15
- [35] A. Benaissa, B. Retiat, B. Cebere, and A. E. Belfedhal, “Tenseal: A library for encrypted tensor operations using homomorphic encryption,” 2021.
- [36] U. von Luxburg, M. Belkin, and O. Bousquet, “Consistency of spectral clustering,” *The Annals of Statistics*, vol. 36, no. 2, pp. 555–586, 04 2008. [Online]. Available: <https://doi.org/10.1214/009053607000000640>
- [37] L. Bottou, F. E. Curtis, and J. Nocedal, “Optimization methods for large-scale machine learning,” *SIAM Review*, vol. 60, no. 2, pp. 223–311, 2018.

Lin Lu (Member, IEEE) received the Ph.D. degree from the School of Computer Science and Technology, Wuhan University of Technology, China. He is currently a lecturer with the college of computer and information technology, China Three Gorges University, China. His research interests include intelligent connected vehicles, big data and federated learning.

Jiaxin Lin is currently pursuing the M.S. degree with the college of computer and information technology, China Three Gorges University, China. His research interests include privacy computation and applied cryptography.

Rodney Donald Hounkanrin is currently pursuing the M.S. degree with the college of computer and information technology, China Three Gorges University, China. His research interests include privacy computation.

Jianting Ning (Senior Member, IEEE) received the Ph.D. degree from the Department of Computer Science and Engineering, Shanghai Jiao Tong University, in 2016. He is currently with the School of Cyber Science and Engineering, Wuhan University, Wuhan, China, and the Faculty of Data Science, City University of Macau, Macau, China. He has published papers in major conferences or journals, such as ACM CCS, NDSS, Asiacrypt, ESORICS, ACSAC, IEEE TIFS, and IEEE TDSC. His research interests include applied cryptography and information security.

SUPPLEMENTARY ANALYSIS OF SECFDCC

This appendix provides the complete mathematical proofs for the theorems presented in the main text, followed by an extended analysis of the privacy guarantees, non-IID robustness mechanisms, and computational complexity of the SecFDCC framework.

A. Proofs of Main Text Theorems

1) Proof of Theorem 1 (Dual-Space Non-IID Robustness):

Proof. Recall the dual-space objective:

$$\alpha \sum_{i,j,k} u_{ijk}^m \left\| \mathbf{x}_i^{(j)} - \mathbf{c}_k \right\|^2 + (1 - \alpha) \sum_{i,j,k} u_{ijk}^m \left\| \hat{\mathbf{h}}_i^{(j)} - \mathbf{e}_k \right\|^2.$$

Let $\mathcal{E}_{\text{orig}}$ and $\mathcal{E}_{\text{emb}}$ denote the optimal clustering errors achievable in the original and embedding spaces, respectively. Suppose, for given memberships and centers, the empirical clustering error in each space is no worse than $\mathcal{E}_{\text{orig}}$ and $\mathcal{E}_{\text{emb}}$:

$$\sum_{i,j,k} u_{ijk}^m \left\| \mathbf{x}_i^{(j)} - \mathbf{c}_k \right\|^2 \leq \mathcal{E}_{\text{orig}}, \quad \sum_{i,j,k} u_{ijk}^m \left\| \hat{\mathbf{h}}_i^{(j)} - \mathbf{e}_k \right\|^2 \leq \mathcal{E}_{\text{emb}}.$$

Aligning all clients' embedding spaces to the global anchor (via Procrustes alignment) induces an additional alignment error per point, denoted ϵ_{align} , measured as the average squared distance between the (aligned) embeddings of anchors and the global embedding centers:

$$\epsilon_{\text{align}} = \frac{1}{MK} \sum_{i=1}^M \sum_{k=1}^K \left\| \hat{\mathbf{h}}_{i,\text{anchor}}^{(k)} - \mathbf{e}_k \right\|^2.$$

By construction, the membership fusion $\alpha U_{\text{orig}} + (1 - \alpha) U_{\text{emb}}$ ensures that the aggregate dual loss cannot exceed the convex combination of separate errors, up to an additive $O(\epsilon_{\text{align}})$ arising from imperfect alignment. Therefore, the practical error achieved by minimizing $\mathcal{L}_{\text{dual}}$ satisfies:

$$\mathcal{E}_{\text{dual}} \leq \alpha \mathcal{E}_{\text{orig}} + (1 - \alpha) \mathcal{E}_{\text{emb}} + O(\epsilon_{\text{align}}).$$

If $\mathcal{E}_{\text{orig}} \approx \mathcal{E}_{\text{emb}}$ and $\alpha = 0.5$, then $\mathcal{E}_{\text{dual}} \leq 0.5 \max(\mathcal{E}_{\text{orig}}, \mathcal{E}_{\text{emb}}) + O(\epsilon_{\text{align}})$, completing the proof. $\square$

2) Proof of Theorem 2 (Alignment Consistency under Non-IID):

Proof. Let $\mathbf{H}_{\text{anchor}}^{(i)} \in \mathbb{R}^{K \times d'}$ be the embedding of the K anchor points at client i , and $\mathcal{C}_{\text{emb}} \in \mathbb{R}^{K \times d'}$ stack the global embedding centers as rows. The orthogonal Procrustes problem finds

$$\mathbf{R}_i^* = \arg \min_{\mathbf{R} \in O(d')} \left\| \mathbf{H}_{\text{anchor}}^{(i)} \mathbf{R} - \mathcal{C}_{\text{emb}} \right\|_F^2.$$

Let $\hat{\mathbf{H}}_{\text{anchor}}^{(i)} = \mathbf{H}_{\text{anchor}}^{(i)} \mathbf{R}_i^*$. By optimality of Procrustes, the error satisfies:

$$\left\| \hat{\mathbf{H}}_{\text{anchor}}^{(i)} - \mathcal{C}_{\text{emb}} \right\|_F^2 \leq \min_{\mathbf{R}} \left\| \mathbf{H}_{\text{anchor}}^{(i)} \mathbf{R} - \mathcal{C}_{\text{emb}} \right\|_F^2.$$

If the embedding at each client is a noisy but unbiased estimator of the global one (ensured by the coupled anchor construction), standard concentration bounds for empirical

spectral embeddings [36] imply the error per anchor decays as $O(\sigma_{\text{noise}}/\sqrt{n_i})$. The residual error from imperfect alignment construction is further bounded by a small constant ϵ_{align} due to regular graph connectivity. Hence:

$$\left\| \hat{\mathbf{H}}_{\text{anchor}}^{(i)} - \mathcal{C}_{\text{emb}} \right\|_F \leq \epsilon_{\text{align}} + O\left(\frac{\sigma_{\text{noise}}}{\sqrt{n_i}}\right).$$

As n_i grows or the federated procedure progresses, both terms shrink. $\square$

3) Proof of Theorem 3 (Convergence of Topology Refinement):

Proof. Let $L_{\text{topo}}(\theta)$ denote the total topology consistency loss as a function of parameters $\theta = (\{\mathbf{c}_k\}_{k=1}^K, \{\mathbf{e}_k\}_{k=1}^K)$. The loss is a sum of:

- Convex quadratic fidelity terms,
- Smooth (mean squared error) topology structure preservation,
- Piecewise linear hinge (separation loss, convex over violated regions).

The overall loss is thus Lipschitz smooth (ignoring the non-differentiable points where separation transitions), with a Lipschitz constant L depending on K and embedding dimension. Applying standard stochastic gradient descent theory for nonconvex, Lipschitz-smooth loss functions (see e.g., [37]), iterates with step size $\eta < 1/L$ satisfy:

$$\min_{0 \leq t < T} \left\| \nabla L_{\text{topo}}(\theta^{(t)}) \right\| \leq \epsilon_{\text{conv}}$$

within $T = O(1/\epsilon_{\text{conv}}^2)$ steps, provided the learning rate is held constant and $\mathcal{L}_{\text{topo}}$ is bounded below. The post-processing normalization (tanh or ℓ_2 projection) is a non-expansive mapping and does not affect convergence guarantees. Early stopping and adaptive schedule further ensure empirical convergence. $\square$

B. Extended Privacy and Security Analysis

We adopt an *honest-but-curious* adversary model where the server follows the protocol but attempts to infer sensitive information. We assume the CKKS scheme is semantically secure under the Ring Learning With Errors (RLWE) assumption.

Theorem 4 (Semantic Security). *Under the RLWE assumption with security parameter $\lambda = 128$, the encrypted statistics $\{\tilde{\mathbf{S}}_{k,\text{orig}}^{(i)}, \tilde{\mathbf{S}}_{k,\text{emb}}^{(i)}, \tilde{n}_k^{(i)}\}$ are computationally indistinguishable from uniformly random polynomials in $\mathcal{R}_Q$ for any PPT adversary $\mathcal{A}$:*

$$\left| \Pr[\mathcal{A}(\tilde{\mathbf{S}}) = 1] - \Pr[\mathcal{A}(r) = 1] \right| \leq \text{negl}(\lambda), \quad (27)$$

where $r \xleftarrow{\$} \mathcal{R}_Q$ is uniformly random.

Proof. This follows directly from the IND-CPA security of the CKKS scheme [34]. Without the secret key $\mathbf{sk}$, distinguishing the ciphertext $(\mathbf{c}_0, \mathbf{c}_1)$ from random noise is equivalent to solving the RLWE problem, which is hard for chosen parameters. $\square$

Theorem 5 (Aggregation Privacy). *The server learns only the global aggregate statistics $\mathbf{S}_k^{\text{global}} = \sum_{i=1}^M \mathbf{S}_{k,*}^{(i)}$ and gains zero additional information about individual contributions beyond what is implied by the aggregate itself.*

Proof. Homomorphic addition ensures $\text{Dec}_{sk}(\sum_i \tilde{\mathbf{S}}_k^{(i)}) = \sum_i \mathbf{S}_k^{(i)} + \mathbf{e}_{\text{agg}}$. The server's view consists solely of indistinguishable ciphertexts (by Theorem 4) and the final decrypted sum. Under the ideal/real world paradigm, a simulator can generate the server's view given only the sum, proving zero information leakage about individual terms [31]. $\square$

Theorem 6 (Accuracy Preservation). *The CKKS decryption error ϵ_{ckks} introduces clustering accuracy degradation bounded by:*

$$|\text{ACC}_{\text{CKKS}} - \text{ACC}_{\text{plaintext}}| \leq O\left(\frac{KM\epsilon_{\text{ckks}}}{\min_k n_k}\right).$$

Proof. The decrypted center is $\hat{\mathbf{c}}_k = (\sum \mathbf{S} + \mathbf{e}_S)/(\sum n + e_n)$. Using Taylor expansion, the perturbation is $\|\hat{\mathbf{c}}_k - \mathbf{c}_k\| \approx O(\|\mathbf{e}\|/n_k)$. Since CKKS noise $\epsilon_{\text{ckks}} \approx 10^{-6}$ and typically $n_k \gg 1$, the displacement is negligible. $\square$

1) Attack Resistance:

- **Membership Inference Attack (MIA):** SecFDCC mitigates MIA via fuzzy clustering ($m = 2$), which spreads a sample's influence, and aggregation over M clients, which dilutes individual contributions. Empirical results show MIA accuracy remains near the random baseline (< 0.55).
- **Model Inversion:** Reconstructing raw data from cluster centers requires solving $\mathbf{S}_k = \sum_j u_{jk}^m \mathbf{x}_j$, a severely underdetermined system ($\text{rank } K \ll n$).
- **Timing Side-Channels:** Constant-time CKKS operations (NTT) yield near-zero correlation ($|r| < 0.1$) between encryption time and data characteristics.

C. Extended Robustness Analysis

We formally analyze robustness against the three specific forms of heterogeneity identified in Section III.

Theorem 7 (Label Skew Robustness). *Assume client i has samples only from a subset $\mathcal{K}_i \subseteq [K]$, and each cluster k is represented in at least ρM clients ($\rho > 0$). Then the global clustering error satisfies:*

$$\mathcal{E}_{\text{SecFDCC}} \leq \max_i \mathcal{E}_i + O\left(\frac{1}{\rho M}\right).$$

Proof. Even if client i lacks samples for cluster k , the anchor embedding mechanism generates a valid embedding for center $\mathbf{c}_k$ based on nearest neighbors in the graph. The weighted Procrustes alignment maintains global coordinate consistency even with missing classes. The global center $\mathbf{c}_k$ is computed by aggregating over the subset of clients possessing cluster k . By the law of large numbers, the estimation error decreases as $O(1/\sqrt{N_k})$, where N_k is the total count across ρM clients. $\square$

Theorem 8 (Feature Shift Robustness). *Spectral embeddings reduce sensitivity to feature shifts by a factor proportional*

to $1/\sqrt{\lambda_{d'}}$, where $\lambda_{d'}$ is the d' -th smallest eigenvalue of the normalized Laplacian.

Proof. Feature shifts act as local affine transformations. The leading eigenvectors of the graph Laplacian capture intrinsic manifold topology, which is invariant to smooth deformations. By Davis-Kahan perturbation theory, high-frequency noise caused by shifts is filtered out, with stability inversely related to the spectral gap $\lambda_{d'+1} - \lambda_{d'}$. $\square$

Theorem 9 (Quantity Imbalance Robustness). *Under severe quantity imbalance ($\max_k n_k / \min_k n_k > 100$), fuzzy clustering with $m = 2$ ensures that minority clusters are not overwhelmed.*

Proof. The fuzzy membership $u_{ijk} \propto (1/d_{ijk}^2)$ ensures that points near a minority center contribute significant weight regardless of the total cluster size n_k . Furthermore, the topology separation loss explicitly prevents the minority center from collapsing into a massive neighbor, ensuring $\mathcal{E}_{\text{minority}}$ remains bounded relative to $\mathcal{E}_{\text{majority}}$. $\square$

D. Convergence Analysis

Theorem 10 (Convergence of SecFDCC under Non-IID). *Under bounded data, sufficient cluster representation (ρM), bounded CKKS error, and a learning rate $\eta < 1/L$, SecFDCC converges to a stationary point of the dual-space objective with rate:*

$$\mathbb{E}[\|\nabla \mathcal{L}_{\text{dual}}^{(t)}\|^2] \leq \frac{\mathcal{L}_{\text{dual}}^{(0)} - \mathcal{L}_{\text{dual}}^*}{T} + O\left(\frac{KM\epsilon_{\text{ckks}}}{\rho M}\right).$$

Proof. The proof combines the descent lemma for the topology optimization (Theorem 3) with the stability bounds of the Procrustes-aligned embeddings (Theorem 2). The noise term $O(KM\epsilon_{\text{ckks}}/\rho M)$ represents the irreducible error floor due to encryption noise and non-IID sampling variance. $\square$

E. Complexity Analysis

1) *Client Computation:* Per round, a client performs:

- **Graph & Embedding:** $O(n_i d'^2)$ via sparse eigensolvers (Lanczos).
- **Procrustes:** $O(K d'^2)$, negligible since $K \ll n_i$.
- **Encryption:** $O(K \cdot N \log N)$, where N is the polynomial degree (e.g., 8192). Using SIMD packing, we batch K statistics into a constant number of ciphertexts.

Total complexity is dominated by the spectral embedding ($O(n_i^2)$ or $O(n_i k)$ depending on implementation) and encryption ($O(1)$ relative to data size).


2) *Server Complexity:*

- **Aggregation:** $O(M \cdot N)$ (homomorphic additions).
- **Topology Refinement:** $O(K^2 d \cdot T_{\text{opt}})$. Since K is small (e.g., 10-50), this is computationally inexpensive (< 1 second).


MORE EXPERIMENTAL RESULTS

F. Non-IID Data Distribution and Visualizations


We visualize SecFDCC’s robustness under Dirichlet non-IID settings ($\alpha = 0.5, M = 20$). In all dataset grids below, the layout is: **Top-Left**: Data distribution; **Top-Right**: Clustering results; **Bottom-Left**: Center evolution; **Bottom-Right**: Ground truth alignment.




(a) Digits Dataset



(b) USPS Dataset



(c) Swiss Roll Dataset



(d) COIL20 Dataset

Fig. 8: Visuals on Swiss Roll and COIL20. Accurate reconstruction of 3D manifolds and global alignment of complex rotational multi-view objects.

a) Summary.: Visual evidence confirms SecFDCC stably maps disjoint client data into a unified space, preserving structural alignment without privacy loss.

Fig. 7: Visuals on Digits and USPS. SecFDCC effectively separates classes despite severe label skew and high-dimensional distribution heterogeneity.